

OraBench - Benchmark Framework for Oracle Database Drivers.

build passing release v1.0.0 release date june github repo or version not found

Table of Contents

- 1. Introduction**
 - 2. Framework Tools**
 - 2.1 Benchmark Configuration**
 - 2.2 Installation**
 - 2.3 Benchmark Operation**
 - 2.4 Benchmark Results**
 - 2.5 Bulk File**
-

1. Introduction

OraBench can be used to determine the performance of different Oracle database drivers under identical conditions. The framework parameters for a benchmark run are stored in a central configuration file.

The currently supported database drivers are:

Driver	Programming Language(s)
cx_Oracle	Python 3
godror	Go
JDBC.jl	Julia
Oracle JDBC	Java & Kotlin
Oracle ODPI-C	C++ (gcc)
Oracle.jl	Julia
oranif	Elixir & Erlang

The following Oracle database versions are provided in a benchmark run via Docker container:

Shortcut	Oracle Database Version
db_18_4_xe	Oracle Database 18c 18.4 (Express Edition) - Linux x86-64
db_19_3_ee	Oracle Database 19c 19.3 - Linux x86-64
db_21_3_ee	Oracle Database 21c 21.3 - Linux x86-64

The results of the benchmark runs are collected in either csv (comma-separated values) or tsv (tab-separated values) files.

2. Framework Tools

2.1 Benchmark Configuration

The benchmark configuration file controls the execution and output of a benchmark run. The default name for the configuration file is `priv/properties/ora_bench.properties`. A detailed description of the configuration options can be found [here](#). For reasons of convenience the following files are generated:

- the configuration file `priv/ora_bench_c.properties` for C++ (gcc),
- the configuration file `priv/ora_bench_erlang.properties` with a corresponding map for Erlang, and
- the configuration file `priv/ora_bench_python.properties` for Python 3.
- the configuration file `priv/ora_bench_toml.properties` for Julia.

All the file names specified here are also part of the configuration file and can be changed if necessary.

2.2 Installation

The easiest way is to download a current release of **OraBench** from the GitHub repository. You can find the necessary link [here](#).

OraBench is tested under [Ubuntu](#).

[Git](#) is needed to download the repository and for compilation the following software components are needed:

- [Elixir](#)
- [Erlang](#)
- [gcc Windows](#) or [gcc Ubuntu](#)
- [Go](#)
- [Gradle Build Tool](#)
- [Java](#), e.g. [the open-source JDK](#)
- [Julia](#)
- [Kotlin](#)
- [Oracle Instant Client](#)
- [Python 3](#)
- [rebar3](#)

For changes to the **OraBench** repository it is best to use an editor (e.g. [Vim](#)) or a suitable IDE. For using the Docker Image based databases in operational mode, [Docker Desktop](#) must also be installed. For the respective software versions, please consult the document [release notes](#).

The whole software environment for the operation and further development of OraBench can be created most easily by using a Docker container (version 1.1.0 from [here](#)).

Alternatively, in an Ubuntu 20.04 based environment, e.g.: in a virtual machine, the two following scripts can be used to install the necessary software:

- `scripts/1.1.0/run_install_4-vm_wsl2_1.sh`

- `scripts/1.1.0/run_install_4-vm_wsl2_2.sh`
 - run `sudo apt update`
 - run `sudo apt install git`
 - run `git clone https://github.com/KonnexionsGmbH/ora_bench` (cloning the **OraBench** repository)
 - run `cd ora_bench/scripts/kxn_dev`
 - run `./run_install_4_vm_wsl2_1.sh`
 - close the Ubuntu shell and reopen it again
 - run `cd ora_bench/scripts/kxn_dev`
 - run `./run_install_4_vm_wsl2_2.sh`

2.3 Benchmark Operation

2.3.1 Script `run_ora_bench`

This script executes the `run_properties_standard` script for each of the databases listed in chapter Introduction with standard properties. At the beginning of the script it is possible to exclude individual databases or drivers from the current benchmark. The run log is stored in the `run_ora_bench.log` file.

2.4 Benchmark Results

In a file defined by the configuration parameters `file.result.delimiter`, `file.result.header` and `file.result.name`, the results of the benchmark run with the actions `benchmark`, `trial` and `query` are stored. In the file directory `priv/statistics` reference statistics files are available per version of **OraBench**.

Excerpts from a sample file can be seen in the following image:

Database	Language	Driver	Duration (ns)
db_21_3_ee	Go gol.17	godror v0.25.3	9133612500
db_19_3_ee	Go gol.17	godror v0.25.3	9307445800
db_18_4_xe	OTP 24, erts-12.0	oranif (Version 0.2.3)	10542000000
db_19_3_ee	OTP 24, erts-12.0	oranif (Version 0.2.3)	10769000000
db_18_4_xe	Go gol.17	godror v0.25.3	12262813800
db_21_3_ee	OTP 24, erts-12.0	oranif (Version 0.2.3)	12393000000
db_21_3_ee	Kotlin 1.5.0	Oracle JDBC (Version 21.3.0.0.0)	18047533700
db_18_4_xe	Kotlin 1.5.0	Oracle JDBC (Version 21.3.0.0.0)	18219792100
db_19_3_ee	Kotlin 1.5.0	Oracle JDBC (Version 21.3.0.0.0)	19543579300
db_21_3_ee	Java 16.0.2	Oracle JDBC (Version 21.3.0.0.0)	20238521000
db_19_3_ee	Java 16.0.2	Oracle JDBC (Version 21.3.0.0.0)	20836544700
db_18_4_xe	Java 16.0.2	Oracle JDBC (Version 21.3.0.0.0)	20859623400
db_18_4_xe	Python 3 3.9.7 (tags/v	Oracle cx_Oracle (Version v8.2.1)	21241982000
db_21_3_ee	Python 3 3.9.7 (tags/v	Oracle cx_Oracle (Version v8.2.1)	21319565000
db_19_3_ee	Python 3 3.9.7 (tags/v	Oracle cx_Oracle (Version v8.2.1)	21817186000
db_21_3_ee	Elixir 1.12.2	oranif (Version 0.2.3)	27548000000
db_18_4_xe	Elixir 1.12.2	oranif (Version 0.2.3)	27805000000
db_19_3_ee	Elixir 1.12.2	oranif (Version 0.2.3)	30766000000

In detail, the following information is available in the result files:

Column	Format	Content
release	alphanumeric	config param <code>benchmark.release</code>

Column	Format	Content
benchmark id	alphanumeric	config param <code>benchmark.id</code>
benchmark comment	alphanumeric	config param <code>benchmark.comment</code>
host name	alphanumeric	config param <code>benchmark.host.name</code>
no. cores	integer	config param <code>benchmark.number.cores</code>
os	alphanumeric	config param <code>benchmark.os</code>
user name	alphanumeric	config param <code>benchmark.user.name</code>
database	alphanumeric	config param <code>benchmark.database</code>
language	alphanumeric	config param <code>benchmark.language</code>
driver	alphanumeric	config param <code>benchmark.driver</code>
trial no.	integer	0 if action equals <code>benchmark</code> , trial no. otherwise
SQL statement	alphanumeric	SQL statement if action equals <code>query</code> , empty otherwise
core multiplier	integer	config param <code>benchmark.core.multiplier</code>
fetch size	integer	config param <code>connection.fetch.size</code>
transaction size	integer	config param <code>benchmark.transaction.size</code>
bulk length	integer	config param <code>file.bulk.length</code>
bulk size	integer	config param <code>file.bulk.size</code>
batch size	integer	config param <code>benchmark.batch.size</code>
action	alphanumeric	one of <code>benchmark</code> , <code>query</code> or <code>trial</code>
start day time	yyyy-mm-dd hh24:mi:ss.fffffffff	current date and time at the start of the action
end day time	yyyy-mm-dd hh24:mi:ss.fffffffff	current date and time at the end of the action
duration (sec)	integer	time difference in seconds between start time and end time of the action
duration (ns)	integer	time difference in nanoseconds between start time and end time of the action

2.5 Bulk File

The bulk file in `csv` or `tsv` format is created in the `run_create_bulk_file` script if it does not already exist. The following configuration parameters are taken into account:

- `file.bulk.delimiter`
- `file.bulk.header`

- `file.bulk.length`
- `file.bulk.name`
- `file.bulk.size`

The data column in the bulk file is randomly generated with a unique key column (MD5 hash code).