

ora_bench - Benchmark Framework for Oracle Database Drivers.

build canceled license not identifiable by github release no releases or repo not found
release date no releases or repo not found github repo or version not found

Table of Contents

- 1. Introduction
 - 2. Framework Tools
 - 3. Coding Pattern
 - 4. Driver Specific Features
 - 5. Execution Variations
 - 6. Reporting
-

1. Introduction

OraBench can be used to determine the performance of different Oracle database drivers under identical conditions. The framework parameters for a benchmark run are stored in a central configuration file.

The currently supported database drivers are:

Driver	Programming Languages
cx_Oracle	Python
godror	Go
JDBC	Java & Kotlin
ODPI	C
oranif	Elixir & Erlang

The following Oracle database versions are provided in a benchmark run via Docker container:

Shortcut	Oracle Database Version
db_18_4_xe	Oracle Database 18c 18.4 (Express Edition) - Linux x86-64
db_19_3_ee	Oracle Database 19c 19.3 - Linux x86-64

The results of the benchmark runs are collected in either csv (comma-separated values) or tsv (tab-separated values) files.

2. Framework Tools

2.1 Benchmark Configuration

The benchmark configuration file controls the execution and output of a benchmark run. The default name for the configuration file is `priv/properties/ora_bench.properties`. A detailed description of the configuration options can be found [here](#). For reasons of convenience the following files are generated:

- the configuration file `priv/ora_bench_c.properties` for C,
- the configuration file `priv/ora_bench_erlang.properties` with a corresponding map for Erlang, and
- the configuration file `priv/ora_bench_python.properties` for Python.

All the file names specified here are also part of the configuration file and can be changed if necessary.

2.2 Benchmark Execution

2.2.1 Locally

2.2.1.1 System Requirements

2.2.1.1.1 Konnexion's Development Image `kxn_dev`

2.2.1.1.2 Windows Platform

- Docker Desktop for Windows from [here](#)
- Make for Windows from [here](#)
- Oracle Instant Client from [here](#)
- Erlang from [here](#)
- Elixir from [here](#)
- Go from [here](#)
- Java SE Development Kit, e.g. Version 11 from [here](#)
- Gradle from [here](#)
- Kotlin from [here](#)
- Python 3 from [here](#)
- rebar3 from [here](#)

2.2.1.2 `run_bench_all_dbs_props_std`

This script executes the `run_properties_standard` script for each of the databases listed in chapter Introduction with standard properties. At the beginning of the script it is possible to exclude individual databases or drivers from the current benchmark. The run log is stored in the `run_bench_all_dbs_props_std.log` file.

2.2.1.3 `run_bench_all_dbs_props_var`

This script executes the `run_properties_variations` script for each of the databases listed in chapter Introduction with variations of properties. At the beginning of the script it is possible to exclude individual databases or drivers from the current benchmark. The run log is stored in the `run_bench_all_dbs_props_var.log` file.

2.2.1.4 `run_bench_<driver>_<programming language>`

The driver and programming language related scripts, such as `run_bench_jdbc` in the `src_java` directory, first execute the insert statements and then the select statements in each trial with the data from the bulk file. The time consumed is captured and recorded in result files.

2.2.2 Travis CI

In Travis CI, the following two environment variables are defined per build for each of the databases listed in chapter Introduction:

- `ORA_BENCH_BENCHMARK_DATABASE`

In each build the following tasks are performed:

1. Installation of Elixir, Erlang, Go, Java, Oracle Instant Client and Python.
2. Creation of the bulk file with the script `run_create_bulk_file`.
3. Execution of the `run_properties_variations` sub-script.
4. Storing the measurement results in the branch `gh-pages`.

2.3 Benchmark Results

In a file defined by the configuration parameters `file.result.delimiter`, `file.result.header` and `file.result.name`, the results of the benchmark run with the actions `benchmark`, `trial` and `query` are stored. If the result file does not yet exist, a new result file is created. Otherwise, the new current results are appended to existing results.

Column	Format	Content
release	alphanumeric	config param <code>benchmark.release</code>
benchmark id	alphanumeric	config param <code>benchmark.id</code>
benchmark comment	alphanumeric	config param <code>benchmark.comment</code>
host name	alphanumeric	config param <code>benchmark.host.name</code>
no. cores	integer	config param <code>benchmark.number.cores</code>
os	alphanumeric	config param <code>benchmark.os</code>
user name	alphanumeric	config param <code>benchmark.user.name</code>
database	alphanumeric	config param <code>benchmark.database</code>
language	alphanumeric	config param <code>benchmark.language</code>

Column	Format	Content
driver	alphanumeric	config param <code>benchmark.driver</code>
trial no.	integer	0 if action equals <code>benchmark</code> , trial no. otherwise
SQL statement	alphanumeric	SQL statement if action equals <code>query</code> , empty otherwise
core multiplier	integer	config param <code>benchmark.core.multiplier</code>
fetch size	integer	config param <code>connection.fetch.size</code>
transaction size	integer	config param <code>benchmark.transaction.size</code>
bulk length	integer	config param <code>file.bulk.length</code>
bulk size	integer	config param <code>file.bulk.size</code>
batch size	integer	config param <code>benchmark.batch.size</code>
action	alphanumeric	one of <code>benchmark</code> , <code>query</code> or <code>trial</code>
start day time	yyyy-mm-dd hh24:mi:ss.fffffffff	current date and time at the start of the action
end day time	yyyy-mm-dd hh24:mi:ss.fffffffff	current date and time at the end of the action
duration (sec)	integer	time difference in seconds between start time and end time of the action
duration (ns)	integer	time difference in nanoseconds between start time and end time of the action

2.4 Bulk File

The bulk file in `csv` or `tsv` format is created in the `run_create_bulk_file` script if it does not already exist. The following configuration parameters are taken into account:

- `file.bulk.delimiter`
- `file.bulk.header`
- `file.bulk.length`
- `file.bulk.name`
- `file.bulk.size`

The data column in the bulk file is randomly generated with a unique key column (MD5 hash code).

3. Coding Patterns

3.1 Benchmark Function (main function)

```
run_benchmark()
```

```
    save the current time as the start of the 'benchmark' action
```

```

        READ the configuration parameters into the memory (config params
`file.configuration.name ...`)
        READ the bulk file data into the partitioned collection
bulk_data_partitions (config param 'file.bulk.name')
        partition key = modulo (ASCII value of 1st byte of key * 256 + ASCII
value of 2nd byte of key,
                                number partitions (config param
'benchmark.number.partitions'))
        Create a separate database connection (without auto commit behaviour) for
each partition

        trial_no = 0
        WHILE trial_no < config_param 'benchmark.trials'
            DO run_benchmark_trial(database connections, trial_no,
bulk_data_partitions)
        ENDWHILE

        partition_no = 0
        WHILE partition_no < config_param 'benchmark.number.partitions'
            close the database connection
        ENDWHILE

        WRITE an entry for the action 'benchmark' in the result file (config param
'file.result.name')

```

3.2 Trial Function

```

run_trial(database connections, trial_no, bulk_data_partitions)
INPUT: the database connections
        the current trial number
        the partitioned bulk data

        save the current time as the start of the 'trial' action

        create the database table (config param 'sql.create')

        IF error
            drop the database table (config param 'sql.drop')
            create the database table (config param 'sql.create')
        ENDIF

        DO run_benchmark_insert(database connections, trial_no,
bulk_data_partitions)
        DO run_benchmark_select(database connections, trial_no,
bulk_data_partitions)

        drop the database table (config param 'sql.drop')

        WRITE an entry for the action 'trial' in the result file (config param
'file.result.name')

```

3.3 Insert Control Function

```

run_insert(database connections, trial_no, bulk_data_partitions)
INPUT: the database connections
       the current trial number
       the partitioned bulk data

       save the current time as the start of the 'query' action

       partition_no = 0
       WHILE partition_no < config_param 'benchmark.number.partitions'
           IF config_param 'benchmark.core.multiplier' = 0
               DO Insert(database connections(partition_no),
bulk_data_partitions(partition_no))
           ELSE
               DO Insert(database connections(partition_no),
bulk_data_partitions(partition_no)) as a thread
           ENDWHILE

       WRITE an entry for the action 'query' in the result file (config param
'file.result.name')

```

3.4 Insert Function

```

insert(database connection, bulk_data_partition)
INPUT: the database connection
       the bulk data partition

       count = 0
       collection batch_collection = empty

       WHILE iterating through the collection bulk_data_partition
           count + 1

           add the SQL statement in config param 'sql.insert' with the current
bulk_data entry to the collection batch_collection
           IF config_param 'benchmark.batch.size' > 0
               IF count modulo config_param 'benchmark.batch.size' = 0
                   execute the SQL statements in the collection batch_collection
                   batch_collection = empty
               ENDIF
           END IF

           IF config_param 'benchmark.transaction.size' > 0 AND count modulo
config_param 'benchmark.transaction.size' = 0
               commit
           ENDIF
       ENDWHILE

```

```

IF collection batch_collection is not empty
    execute the SQL statements in the collection batch_collection
ENDIF

commit

```

3.5 Select Control Function

```

run_select(database connections, trial_no, bulk_data_partitions)
INPUT: the database connections
       the current trial number
       the partitioned bulk data

save the current time as the start of the 'query' action

partition_no = 0
WHILE partition_no < config_param 'benchmark.number.partitions'
    IF config_param 'benchmark.core.multiplier' = 0
        DO Select(database connections(partition_no),
bulk_data_partitions(partition_no, partition_no)
    ELSE
        DO Select(database connections(partition_no),
bulk_data_partitions(partition_no, partition_no) as a thread
    ENDWHILE

WRITE an entry for the action 'query' in the result file (config param
'file.result.name')

```

3.6 Select Function

```

run_select(database connection, bulk_data_partition, partition_no)
INPUT: the database connection
       the bulk data partition
       the current partition number

save the current time as the start of the 'query' action

count = 0

execute the SQL statement in config param 'sql.select'

WHILE iterating through the result set
    count + 1
ENDWHILE

IF NOT count = size(bulk_data_partition)
    display an error message
ENDIF

```

4. Driver Specific Features

4.1 cx_Oracle and Python

- the following data in the configuration parameters is determined at runtime:
 - cx_Oracle version (`benchmark.driver`) and
 - Python version (`benchmark.language`).
- all configuration parameters are managed by the program OraBench.java and made available in a suitable file (`file.configuration.name.python`)
- Python uses for batch operations the `executemany` method of the `cursor` class for the operation `INSERT`
- the value fetch size (`connection.fetch.size`) is not used because the operation `SELECT` uses the operation `Cursor.fetchall()`

4.2 JDBC and Java

- the following data in the configuration parameters is determined at runtime:
 - JDBC version (`benchmark.driver`),
 - benchmark identifier (`benchmark.id`),
 - host name (`benchmark.host.name`),
 - number of cores (`benchmark.number.cores`),
 - JRE version (`benchmark.language`),
 - operating system environment (`benchmark.os`),
 - user name (`benchmark.user.name`) and
 - SQL create statement (`sql.create`).
- the Java source code is compiled with the help of Gradle
- Java uses the `PreparedStatement` class for the operations `INSERT` and `SELECT`
- Java uses for batch operations the `executeBatch` method of the `PreparedStatement` class for the operation `INSERT`

4.3 JDBC and Kotlin

- the following data in the configuration parameters is determined at runtime:
 - Kotlin version (`benchmark.driver`),
 - benchmark identifier (`benchmark.id`),
 - host name (`benchmark.host.name`),
 - number of cores (`benchmark.number.cores`),
 - JRE version (`benchmark.language`),
 - operating system environment (`benchmark.os`),
 - user name (`benchmark.user.name`) and
 - SQL create statement (`sql.create`).
- the Kotlin source code is compiled with the help of Gradle
- Kotlin uses the `PreparedStatement` class for the operations `INSERT` and `SELECT`
- Kotlin uses for batch operations the `executeBatch` method of the `PreparedStatement` class for the operation `INSERT`

4.4 ODPI and C

- the following data in the configuration parameters is determined at runtime:
 - ODPI version (`benchmark.driver`) and
 - C version (`benchmark.language`).
- all configuration parameters are managed by the program OraBench.java and made available in a suitable file (`file.configuration.name.c`)

4.5 oranif and Elixir

- the following data in the configuration parameters is determined at runtime:
 - oranif version (`benchmark.driver`) and
 - Elixir version (`benchmark.language`).

4.6 oranif and Erlang

- the following data in the configuration parameters is determined at runtime:
 - oranif version (`benchmark.driver`) and
 - Erlang version (`benchmark.language`).
- all configuration parameters are managed by the program OraBench.java and made available in a suitable file (`file.configuration.name.erlang`)

5. Execution Variations

5.1 Ubuntu 20.04 LTS (including VMware)

- **Requirements:**
 - Ubuntu 20.04 installed directly or via VMware
 - run `sudo apt update`
 - run `sudo apt install dos2unix git`
 - add the following lines to `.bash_profile`:

```
if [ -f ~/.bashrc ]; then
    source ~/.bashrc
fi
```

- run `export DOCKER_USERNAME=\<user name\>`
- run `export DOCKER_PASSWORD=\<password\>`
- run `git clone https://github.com/KonnexionsGmbH/ora_bench` (cloning the ora_bench repository)
- run `cd ora_bench`
- run `./scripts/run_prep_bash_scripts.sh` (preparing the shell scripts)

- run `./scripts/run_install_4_ubuntu_20.04_vm_wsl2.sh` (setting up the WSL2 environment)
- close the Ubuntu shell and reopen it again
- run `cd ora_bench`
- **Execution:** run `./run_bench_all_dbs_props_[std|var].sh`

5.2 Ubuntu 20.04 LTS and `kxn_dev` Image

- **Requirements:**
 - pull the `kxn_dev` image from DockerHub: `docker pull konnexionsgmbh/kxn_dev:latest`
 - create an appropriate container: `docker run -it --name kxn_dev -v /var/run/docker.sock:/var/run/docker.sock konnexionsgmbh/kxn_dev:latest bash`
 - run `export DOCKER_USERNAME=<user name>`
 - run `export DOCKER_PASSWORD=<password>`
 - run `git clone https://github.com/KonnexionsGmbH/ora_bench` (cloning the `ora_bench` repository)
 - run `cd ora_bench`
 - run `./scripts/run_prep_bash_scripts.sh` (preparing the shell scripts)
 - run `gradle copyJarToLib`
- **Execution:** `./run_ora_bench.sh`
- **Issues:**
 - Trino Distributed Query Engine and Microsoft SQL Connector

5.3 Ubuntu 20.04 LTS and Windows Subsystem Linux 2

- **Requirements:**
 - install Ubuntu 20.04 from Microsoft Marketplace
 - run `sudo apt update`
 - run `sudo apt install dos2unix`
 - add the following lines to `.bash_profile`:

```
if [ -f ~/.bashrc ]; then
    source ~/.bashrc
fi
```

- activate the `WSL INTEGRATION` for Ubuntu 20.04 in Docker
- **Requirements (continued):**
 - run `export DOCKER_USERNAME=<user name>`

- run `export DOCKER_PASSWORD=\<password\>`
- run `git clone https://github.com/KonnexionsGmbH/ora_bench` (cloning the ora_bench repository)
- run `cd ora_bench`
- run `./scripts/run_prep_bash_scripts.sh` (preparing the shell scripts)
- run `./scripts/run_install_4_ubuntu_20.04_vm_wsl2.sh` (setting up the WSL2 environment)
- close the Ubuntu shell and reopen it again
- run `cd ora_bench`
- run `gradle copyJarToLib`
- **Execution:** run `./run_ora_bench.sh`
- **Issues:**
 - Trino Distributed Query Engine and Microsoft SQL Connector
 - YugabyteDB and Docker image

5.4 Windows 10 Pro

- **Requirements:**
 - run `set DOCKER_USERNAME=\<user name\>`
 - run `set DOCKER_PASSWORD=\<password\>`
 - run `git clone https://github.com/KonnexionsGmbH/ora_bench` (cloning the ora_bench repository)
 - run `cd ora_bench`
- **Execution:** run `run_ora_bench.bat`
- **Issues:**
 - Trino Distributed Query Engine and Microsoft SQL Connector
 - YugabyteDB and Docker image

6. Reporting

[see here](#)